

Solving Large-Scale, Sparse, Constrained Nonlinear Optimization Problems on GPUs

Pivoting-Free Interior Point Methods

Sungho Shin
shin.mit.edu

Massachusetts Institute of Technology

2025 INFORMS Annual Meeting
October 27, 2025

Who are we?

<https://madsuite.org/>



- ▶ Alexis Montois @ Argonne National Lab
- ▶ François Pacaud @ MINES Paris-PSL (an ANL alumnus)
- ▶ Sungho Shin @ MIT (an ANL alumnus)
- ▶ Mihai Aniteșcu @ Argonne National Lab
- ▶ and friends... Michael Saunders, Dominic Orban, Armin Nurkanović, Anton Pozharskiy, Jean-Baptiste Caillau, ...

GPU-based optimization

- ▶ What is GPU? —many-core accelerator with high memory bandwidth
- ▶ Traditionally used for graphics, now widely used for ML/AI workloads
- ▶ Classical optimization (mathematical program; not ML training problems) has been thought to be **unsuitable** for GPUs
- ▶ Many state-of-the-art solvers (on CPUs) rely on **sparse direct factorizations**, which has been considered **difficult** to implement efficiently on GPUs
- ▶ This status quo is changing—the focus of this talk

GPU-based optimization: current landscape

- ▶ **Factorization-free methods** have been the dominant approach
 - ▶ PDLP: Primal-dual first-order methods (Applegate et al., 2022; Lu et al., 2025)
 - ▶ PCG-based algorithms for convex QPs (Schubiger et al., 2020)
 - ▶ ALM + IPM + PCG for nonlinear problems (Cao et al., 2016)
- ▶ Factorization-free methods only rely on **GPU-friendly kernels**—`spmv`, `axpy`, `map`, `mapreduce`—makes them easy to implement on GPUs
- ▶ Scale to extremely large problems **but achieving high precision is difficult**
- ▶ Can we implement **second-order methods** with **sparse direct solvers** on GPUs?

David Applegate et al. (2022). *Practical Large-Scale Linear Programming Using Primal-Dual Hybrid Gradient*. [arXiv: 2106.04756 \[math\]](#)

Haihao Lu et al. (2025). *cuPDL+: A Further Enhanced GPU-Based First-Order Solver for Linear Programming*. [arXiv: 2507.14051 \[math\]](#)

Michel Schubiger et al. (2020). "GPU Acceleration of ADMM for Large-Scale Quadratic Programming". In: *Journal of Parallel and Distributed Computing*

Yankai Cao et al. (2016). "An Augmented Lagrangian Interior-Point Approach for Large-Scale NLP Problems on Graphics Processing Units". In: *Computers & Chemical Engineering*

Our problem setting

Large-scale, sparse, constrained nonlinear programs

- ▶ We consider:

$$\begin{aligned} \min_{x^b \leq x \leq x^\#} \quad & f(x) \\ \text{s.t.} \quad & g^b \leq g(x) \leq g^\#, \end{aligned}$$

... but will particularly focus on:

- ▶ $f : \mathbb{R}^n \rightarrow \mathbb{R}$ and $g : \mathbb{R}^n \rightarrow \mathbb{R}^m$ twice continuously differentiable
- ▶ n and m are large (millions or more)
- ▶ $\nabla g(x)$ and $\nabla_{xx}^2 L(x, \lambda)$ are **sparse** (tens of nonzeros per row)— $L(\cdot, \cdot)$ is the Lagrangian
- ▶ Generally **nonconvex**—we only seek **local solutions**
- ▶ Most important instances (e.g., energy systems, optimal control) fall into this category
- ▶ Most existing solvers (e.g., Ipopt, KNITRO) are optimized for such problems

Our problem setting

Large-scale, sparse, constrained nonlinear programs

- ▶ WLOG (via $a = b \iff a \geq b, b \geq a$), we consider:

$$\begin{aligned} \min_{x \in \mathbb{R}^n} \quad & f(x) \\ \text{s.t.} \quad & g(x) \geq 0, \end{aligned}$$

... but will particularly focus on:

- ▶ $f : \mathbb{R}^n \rightarrow \mathbb{R}$ and $g : \mathbb{R}^n \rightarrow \mathbb{R}^m$ twice continuously differentiable
- ▶ n and m are large (millions or more)
- ▶ $\nabla g(x)$ and $\nabla_{xx}^2 L(x, \lambda)$ are **sparse** (tens of nonzeros per row)— $L(\cdot, \cdot)$ is the Lagrangian
- ▶ Generally **nonconvex**—we only seek **local solutions**
- ▶ Most important instances (e.g., energy systems, optimal control) fall into this category
- ▶ Most existing solvers (e.g., Ipopt, KNITRO) are optimized for such problems

NLP software stack—either on CPUs or GPUs

(1) algebraic modeling system, (2) NLP solver, and (3) sparse, direct linear solver

- ▶ Algebraic modeling systems (e.g., JuMP, AMPL, CasADi) generate **callback** functions— $f(\cdot)$, $g(\cdot)$, $\nabla f(\cdot)$, $\nabla g(\cdot)$, $\nabla_{xx}^2 L(\cdot, \cdot)$ —for NLP solvers
- ▶ NLP solvers (e.g., Ipopt, KNITRO) form **KKT systems** to find step direction:

$$\underbrace{\begin{bmatrix} \nabla_{xx}^2 L(x, s, \lambda) + \delta_p I & & \nabla g(x)^\top \\ & S^{-1} \Lambda & -I \\ \nabla g(x) & -I & -\delta_d I \end{bmatrix}}_{\text{KKT matrix}} \underbrace{\begin{bmatrix} d_x \\ d_s \\ -d_\lambda \end{bmatrix}}_{\text{step direction}} = - \underbrace{\begin{bmatrix} \nabla f(x) - \nabla g(x)^\top \lambda \\ \Lambda e - \mu S^{-1} e \\ g(x) - s \end{bmatrix}}_{\text{residual to KKT conditions}},$$

where s is the slack variable, $S = \text{diag}(s)$, $\Lambda = \text{diag}(\lambda)$, and $\delta_p, \delta_d \geq 0$ are primal-dual regularization parameters

- ▶ **Sparse, direct linear solvers** (e.g., MA27, PARDISO) factorize and solve KKT systems; iterative method is typically not an option due to ill-conditioning

Full NLP software stack on GPU—is it possible?

(1) algebraic modeling system, (2) NLP solver, and (3) sparse, direct linear solver

- ▶ Partially porting on GPU (e.g., AD or linear algebra only) is not sufficient—will be bottlenecked by CPU–GPU communication. We need **all three components on GPU**
- ▶ Porting **algebraic modeling system** is relatively straightforward. ExaModels.jl provides GPU-compatible modeling capabilities (Shin et al., 2024)—not the focus today
- ▶ Porting **direct linear solver** is more non-trivial; as of now, **there is no drop-in replacement of MA27 on GPUs**

Our approach

Adapt **interior-point method for NLPs** so that it can work with **GPU direct linear solvers** with (relatively) limited capabilities

Sparse direct solvers within interior-point methods

LBL^\top factorization

$$\begin{bmatrix} \nabla_{xx}^2 L(x, s, \lambda) + \delta_p I & & \nabla g(x)^\top \\ & S^{-1} \Lambda & -I \\ \nabla g(x) & -I & -\delta_d I \end{bmatrix} = P \times Q \times L \times B \times L^\top \times Q^\top \times P^\top$$

- ▶ P : Fill-in-reducing (and parallel-friendly) re-ordering (re-used across IPM iterations)
- ▶ Q : Pivoting (computed on-the-fly during factorization)
- ▶ LBL^\top : Factorization of permuted matrix
- ▶ Additionally, pivots may be **perturbed**, which must be corrected via **iterative refinement**
- ▶ Factorization reveals inertia—number of $+$, 0 , and $-$ eigenvalues; primal-dual regularization (δ_p, δ_d) is adjusted to achieve correct inertia (Wächter et al., 2006); this can be merged with pivot perturbation (Nocedal et al., 2006)

Pivoting-free IPM

Montoison, Pacaud, Shin, et al., 2025

- ▶ Numerical pivoting is **difficult** to implement on GPUs (Świrydowicz et al., 2022) because row swapping can potentially **destroy parallelism**
- ▶ Existing GPU solvers (e.g., cuDSS) **don't have adequately robust pivoting strategies**; currently only have partial pivoting and pivot perturbation (NVIDIA, 2024)
- ▶ Factorization may fail without **numerical pivoting**, unless
 - ▶ **Symmetric quasi-definite** (SQD): factorization exists (Vanderbei, 1995)
 - ▶ **Symmetric positive definite** (SPD): factorization exists and numerically stable

However, **KKT systems for NLPs are generally indefinite**

- ▶ When pivoting is not required, **GPU sparse direct solvers can be very efficient**

Can we make IPM pivoting-free?

Alexis Montoison, François Pacaud, Sungho Shin, et al. (2025). *GPU Implementation of Second-Order Linear and Nonlinear Programming Solvers*. arXiv: 2508.16094 [math]

Robert J. Vanderbei (1995). "Symmetric Quasidefinite Matrices". In: *SIAM Journal on Optimization*

Kasia Świrydowicz et al. (2022). "Linear Solvers for Power Grid Optimization Problems: A Review of GPU-accelerated Linear Solvers". In: *Parallel Computing*

NVIDIA (2024). *NVIDIA cuDSS (Preview): A High-Performance CUDA Library for Direct Sparse Solvers* — NVIDIA cuDSS Documentation.

<https://docs.nvidia.com/cuda/cudss/index.html>

Pivoting-free IPM

Montoison, Pacaud, Shin, et al., 2025

- ▶ Numerical pivoting is **difficult** to implement on GPUs (Świrydowicz et al., 2022) because row swapping can potentially **destroy parallelism**
- ▶ Existing GPU solvers (e.g., cuDSS) **don't have adequately robust pivoting strategies**; currently only have partial pivoting and pivot perturbation (NVIDIA, 2024)
- ▶ Factorization may fail without **numerical pivoting**, unless
 - ▶ **Symmetric quasi-definite** (SQD): factorization exists (Vanderbei, 1995)
 - ▶ **Symmetric positive definite** (SPD): factorization exists and numerically stable

However, **KKT systems for NLPs are generally indefinite**

- ▶ When pivoting is not required, **GPU sparse direct solvers can be very efficient**

(For now) can we reformulate KKT systems to be SPD or SQD?

Alexis Montoison, François Pacaud, Sungho Shin, et al. (2025). *GPU Implementation of Second-Order Linear and Nonlinear Programming Solvers*. arXiv: 2508.16094 [math]

Robert J. Vanderbei (1995). "Symmetric Quasidefinite Matrices". In: *SIAM Journal on Optimization*

Kasia Świrydowicz et al. (2022). "Linear Solvers for Power Grid Optimization Problems: A Review of GPU-accelerated Linear Solvers". In: *Parallel Computing*

NVIDIA (2024). *NVIDIA cuDSS (Preview): A High-Performance CUDA Library for Direct Sparse Solvers* — NVIDIA cuDSS Documentation.

<https://docs.nvidia.com/cuda/cudss/index.html>

Option 1: Lifted KKT system

Shin et al., 2024

- First, relax equalities by small enough τ (the same as solver tolerances):

$$-\tau \leq h(x) \leq \tau$$

- **Condense** the KKT system by eliminating slack/dual variables:

$$\begin{bmatrix} \nabla_{xx}^2 L(x, s, \lambda) + \delta_p I & & \nabla g(x)^\top \\ & S^{-1} \Lambda & -I \\ \nabla g(x) & -I & -\delta_d I \end{bmatrix} \begin{bmatrix} d_x \\ d_s \\ -d_\lambda \end{bmatrix} = - \begin{bmatrix} \nabla f(x) - \nabla g(x)^\top \lambda \\ \Lambda e - \mu S^{-1} e \\ g(x) - s \end{bmatrix}$$
$$\iff \underbrace{\left(\nabla_{xx}^2 L(x, s, \lambda) + \delta_p I + \nabla g(x)^\top (\delta_d I + \Lambda^{-1} S)^{-1} \nabla g(x) \right)}_{\text{SPD iff the original system has correct inertia}} d_x = -r_p$$

(and some easy-to-solve equations)

- Enables numerically stable factorization **without pivoting or excessive regularization**

Sungho Shin et al. (2024). "Accelerating Optimal Power Flow with GPUs: SIMD Abstraction of Nonlinear Programs and Condensed-Space Interior-Point Methods". In: *Electric Power Systems Research*

Option 1: Lifted KKT—numerical results (AC OPF)

Montoison, Pacaud, Shin, et al., 2025

	Tol	Solver	Small $\text{nnz} < 2^{18}$		Medium $2^{18} \leq \text{nnz} < 2^{20}$		Large $2^{20} \leq \text{nnz}$		Total	
			Solved	Time	Solved	Time	Solved	Time	Solved	Time
OPF	10^{-4}	MadNLP (GPU)	31	0.4166	24	2.6380	11	3.7040	66	1.6979
		Ipopt (CPU)	31	0.3970	24	5.0697	11	38.5053	66	5.3817
	10^{-8}	MadNLP (GPU)	30	2.5037	24	4.6016	10	12.8040	64	4.6228
		Ipopt (CPU)	31	0.5100	24	5.4292	11	37.7818	66	5.5541

- **Benchmark set**: AC optimal power flow (Babaeinejadsarookolae et al., 2021)
- **Hardware**: NVIDIA GV 100 GPU / Intel Xeon Gold 6130 CPU
- **Time**: reported as $\text{SGM10} := \prod_{i=1}^n (t_i + 10))^{1/n} - 10$, max wall time 900s

Observation

GPU solver is most effective for large-scale problems for moderate accuracy (10^{-4})

Option 1: Lifted KKT—numerical results (COPS)

Montoison, Pacaud, Shin, et al., 2025

	Tol	Solver	Small $\text{nnz} < 2^{18}$		Medium $2^{18} \leq \text{nnz} < 2^{20}$		Large $2^{20} \leq \text{nnz}$		Total	
			Solved	Time	Solved	Time	Solved	Time	Solved	Time
COPS	10^{-4}	MadNLP (GPU)	13	0.8665	15	4.8665	16	3.8194	44	3.2314
		Ipopt (CPU)	13	5.2315	15	15.9701	15	45.8411	43	19.2243
	10^{-8}	MadNLP (GPU)	13	0.8575	16	1.5572	16	8.3549	45	3.3797
		Ipopt (CPU)	13	5.9413	15	17.6758	15	40.8639	43	19.2999

- **Benchmark set:** COPS Benchmark (Dolan et al., 2001)
- **Hardware:** NVIDIA GV 100 GPU / Intel Xeon Gold 6130 CPU
- **Time:** reported as $\text{SGM10} := \prod_{i=1}^n (t_i + 10))^{1/n} - 10$, max wall time 900s

Observation

GPU solver is most effective for large-scale problems for moderate accuracy (10^{-4})

Option 2: Hybrid KKT system

Golub et al., 2003; Pacaud et al., 2024; Regev et al., 2023

- We keep the equality constraints $h(x) = 0$ and eliminate the inequalities (slack/dual):

$$\begin{bmatrix} K_{\text{cond}} + \delta_p I & \nabla h(x)^\top \\ \nabla h(x) & -\delta_d I \end{bmatrix} \begin{bmatrix} d_x \\ -d_\lambda \end{bmatrix} = - \begin{bmatrix} r_x \\ r_\lambda \end{bmatrix}$$

$\xrightarrow[\text{(w/o changing solution)}]{\text{Regularization}}$

$$\begin{bmatrix} K_{\text{cond}} + \delta_p I + \rho \nabla h(x)^\top \nabla h(x) & \nabla h(x)^\top \\ \nabla h(x) & -\delta_d I \end{bmatrix} \begin{bmatrix} d_x \\ -d_\lambda \end{bmatrix} = - \begin{bmatrix} r'_x \\ r_\lambda \end{bmatrix}$$

Then, the Schur complement can be solved with an iterative method (e.g., CG)

$$\underbrace{\left(\nabla h(x)^\top \left(\overbrace{K_{\text{cond}} + \delta_p I + \rho \nabla h(x)^\top \nabla h(x)}^{\text{SPD for sufficiently large } \rho \text{ iff correct inertia}} \right)^{-1} \nabla h(x) + \delta_d I \right)}_{\text{converges to } I \text{ as } \rho \rightarrow \infty} d_\lambda = r'_\lambda$$

Gene H. Golub et al. (2003). "On Solving Block-Structured Indefinite Linear Systems". In: *SIAM Journal on Scientific Computing*

Shaked Regev et al. (2023). "HyKKT: A Hybrid Direct-Iterative Method for Solving KKT Linear Systems". In: *Optimization Methods and Software*

François Pacaud et al. (2024). *Condensed-Space Methods for Nonlinear Programming on GPUs*. [arXiv: 2405.14236 \[math\]](https://arxiv.org/abs/2405.14236)

Options 3, 4, 5, ...

- ▶ There can be lots of other ways to implement pivoting-free solvers
- ▶ **MadIPM**: Solver for convex quadratic programs; uses regularization to avoid pivoting (Montoison, Pacaud, Shin, et al., 2025)
- ▶ **MadNCL**: ALM-based algorithm—ALM is more **GPU-friendly** due to the quadratic penalty term. Uses nonlinearly-constrained Lagrangian (NCL) formulation and applies condensation techniques (Montoison, Pacaud, Saunders, et al., 2025)
- ▶ Directly solving augmented systems without pivoting? —currently under investigation

Concluding remarks

- ▶ It is an exciting time to develop optimization solvers!
- ▶ Existing GPU solvers have focused on **factorization-free** (e.g. first-order) methods, but
GPU optimization solver \supsetneq Factorization-free solver.
- ▶ With **GPU direct solvers**, **second-order methods** on GPUs has now become possible
- ▶ To harness these GPU direct solvers, we need **pivoting-free algorithms**
- ▶ We presented **Lifted KKT** and **Hybrid KKT**, but there can be many others

Check out our project page!

<https://madsuite.org/>

References

- Applegate, David et al. (2022). *Practical Large-Scale Linear Programming Using Primal-Dual Hybrid Gradient*. arXiv: 2106.04756 [math].
- Babaeinejadsarookolaei, Sogol et al. (2021). *The Power Grid Library for Benchmarking AC Optimal Power Flow Algorithms*. arXiv: 1908.02788 [math].
- Cao, Yankai et al. (2016). "An Augmented Lagrangian Interior-Point Approach for Large-Scale NLP Problems on Graphics Processing Units". In: *Computers & Chemical Engineering* 85, pp. 76–83.
- Dolan, E. D. et al. (2001). *Benchmarking Optimization Software with COPS*. Tech. rep. ANL/MCS-TM-246. Argonne National Lab., IL (US).
- Golub, Gene H. et al. (2003). "On Solving Block-Structured Indefinite Linear Systems". In: *SIAM Journal on Scientific Computing* 24.6, pp. 2076–2092.
- Lu, Haihao et al. (2025). *cuPDLP+: A Further Enhanced GPU-Based First-Order Solver for Linear Programming*. arXiv: 2507.14051 [math].
- Montoison, Alexis, François Pacaud, Michael Saunders, et al. (2025). *MadNCL: A GPU Implementation of Algorithm NCL for Large-Scale, Degenerate Nonlinear Programs*. arXiv: 2510.05885 [math].
- Montoison, Alexis, François Pacaud, Sungho Shin, et al. (2025). *GPU Implementation of Second-Order Linear and Nonlinear Programming Solvers*. arXiv: 2508.16094 [math].
- Nocedal, Jorge et al. (2006). *Numerical Optimization*. 2nd ed. Springer Series in Operations Research. New York: Springer.
- NVIDIA (2024). *NVIDIA cuDSS (Preview): A High-Performance CUDA Library for Direct Sparse Solvers — NVIDIA cuDSS Documentation*. <https://docs.nvidia.com/cuda/cudss/index.html>.
- Pacaud, François et al. (2024). *Condensed-Space Methods for Nonlinear Programming on GPUs*. arXiv: 2405.14236 [math].
- Regev, Shaked et al. (2023). "HyKKT: A Hybrid Direct-Iterative Method for Solving KKT Linear Systems". In: *Optimization Methods and Software* 38.2, pp. 332–355.
- Schubiger, Michel et al. (2020). "GPU Acceleration of ADMM for Large-Scale Quadratic Programming". In: *Journal of Parallel and Distributed Computing* 144, pp. 55–67.
- Shin, Sungho et al. (2024). "Accelerating Optimal Power Flow with GPUs: SIMD Abstraction of Nonlinear Programs and Condensed-Space Interior-Point Methods". In: *Electric Power Systems Research* 236, p. 110651.
- Świrydowicz, Kasia et al. (2022). "Linear Solvers for Power Grid Optimization Problems: A Review of GPU-accelerated Linear Solvers". In: *Parallel Computing* 111, p. 102870.
- Vanderbei, Robert J. (1995). "Symmetric Quasidefinite Matrices". In: *SIAM Journal on Optimization* 5.1, pp. 100–113.
- Wächter, Andreas et al. (2006). "On the Implementation of an Interior-Point Filter Line-Search Algorithm for Large-Scale Nonlinear Programming". In: *Mathematical Programming* 106.1, pp. 25–57.